

Genie: A new, fast, and outlier-resistant hierarchical clustering algorithm

Marek Gagolewski^{a,b,*}, Maciej Bartoszek^{b,c}, Anna Cena^{a,c}

^a*Systems Research Institute, Polish Academy of Sciences
ul. Newelska 6, 01-447 Warsaw, Poland*

^b*Faculty of Mathematics and Information Science, Warsaw University of Technology
ul. Koszykowa 75, 00-662 Warsaw, Poland*

^c*International PhD Studies Program,
Institute of Computer Science, Polish Academy of Sciences*

Abstract

The time needed to apply a hierarchical clustering algorithm is most often dominated by the number of computations of a pairwise dissimilarity measure. Such a constraint, for larger data sets, puts at a disadvantage the use of all the classical linkage criteria but the single linkage one. However, it is known that the single linkage clustering algorithm is very sensitive to outliers, produces highly skewed dendrograms, and therefore usually does not reflect the true underlying data structure – unless the clusters are well-separated. To overcome its limitations, we propose a new hierarchical clustering linkage criterion called Genie. Namely, our algorithm links two clusters in such a way that a chosen economic inequity measure (e.g., the Gini- or Bonferroni-index) of the cluster sizes does not drastically increase above a given threshold. The presented benchmarks indicate a high practical usefulness of the introduced method: it most often outperforms the Ward or average linkage in terms of the clustering quality while retaining the single linkage’s speed. The Genie algorithm is easily parallelizable and thus may be run on multiple threads to speed up its execution even further. Its memory overhead is small: there is no need to precompute the complete distance matrix to perform the computations in order to obtain a desired clustering. It can be applied on arbitrary spaces equipped with a dissimilarity measure, e.g., on real vectors, DNA or protein sequences, images, rankings, informetric data, etc. A reference implementation of the algorithm has been included in the open source `genie` package for R. [See <https://genieclust.gagolewski.com> for a new implementation – available for both R and Python.]

Please cite this paper as: Gagolewski M., Bartoszek M., Cena A., Genie: A new, fast, and outlier-resistant hierarchical clustering algorithm, *Information Sciences* **363**, 2016, pp. 8–23, doi:10.1016/j.ins.2016.05.003.

*Corresponding author; Email: marek.gagolewski@pw.edu.pl

Keywords: hierarchical clustering, single linkage, inequity measures, Gini-index

1. Introduction

Cluster analysis, compare [31], is one of the most commonly applied unsupervised machine learning techniques. Its aim is to automatically discover an underlying structure of a given data set $\mathcal{X} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}\}$ in a form of a partition of its elements: disjoint and nonempty subsets are determined in such a way that observations within each group are “similar” and entities in distinct clusters “differ” as much as possible from each other. This contribution focuses on classical hierarchical clustering algorithms [10, 14] which determine a sequence of nested partitions, i.e., a whole hierarchy of data set subdivisions that may be cut at an arbitrary level and may be computed based on a pairwise dissimilarity measure $\mathfrak{d} : \mathcal{X} \times \mathcal{X} \rightarrow [0, \infty]$ that fulfills *very mild assumptions*: (a) \mathfrak{d} is symmetric, i.e., $\mathfrak{d}(\mathbf{x}, \mathbf{y}) = \mathfrak{d}(\mathbf{y}, \mathbf{x})$ and (b) $(\mathbf{x} = \mathbf{y}) \implies \mathfrak{d}(\mathbf{x}, \mathbf{y}) = 0$ for any $\mathbf{x}, \mathbf{y} \in \mathcal{X}$. This group of clustering methods is often opposed to – among others – partitional schemes which require the number of output clusters to be set up in advance: these include the k -means, k -medians, k -modes, or k -medoids algorithms [34, 40, 56, 59] and fuzzy clustering schemes [6, 46–48], or the BIRCH (balanced iterative reducing and clustering using hierarchies) method [60] that works on real-valued vectors only.

In the large and big data era, one often is faced with the need to cluster data sets of considerable sizes, compare, e.g., [30]. If the $(\mathcal{X}, \mathfrak{d})$ space is “complex”, we observe that the run-times of hierarchical clustering algorithms are dominated by the cost of pairwise distance (dissimilarity measure) computations. This is the case of, e.g., DNA sequences or ranking clustering, where elements in \mathcal{X} are encoded as integer vectors, often of considerable lengths. Here, one often relies on such computationally demanding metrics as the Levenshtein or the Kendall one. Similar issues appear in numerous other application domains, like pattern recognition, knowledge discovery, image processing, bibliometrics, complex networks, text mining, or error correction, compare [10, 15, 16, 18, 29].

In order to achieve greater speed-up, most hierarchical clustering algorithms are applied on a precomputed distance matrix, $(d_{i,j})_{i < j}$, $d_{i,j} = \mathfrak{d}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$, so as to avoid determining the dissimilarity measure for each unique (unordered) pair more than once. This, however, drastically limits the size of an input data set that may be processed. Assuming that $d_{i,j}$ is represented with the 64-bit floating-point (IEEE double precision) type, already the case of $n = 100,000$ objects is way beyond the limits of personal computers popular nowadays: the sole distance matrix would occupy ca. 40GB of available RAM. Thus, for “complex” data domains, we must require that the number of calls to \mathfrak{d} is kept as small as possible. This practically disqualifies all popular hierarchical clustering approaches other than the *single linkage* criterion, for which there is a fast $O(n^2)$ -time and $O(n)$ -space al-

gorithm, see [41, 42, 56], that requires each $d_{i,j}$, $i < j$, to be computed exactly once, i.e., there are precisely $(n^2 - n)/2$ total calls to \mathfrak{d} .

Nevertheless, the single linkage criterion is not eagerly used by practitioners. This is because it is highly sensitive to observations laying far away from clusters' boundaries (e.g., outliers). Because of that, it tends to produce highly skewed dendrograms: at its higher levels one often finds a single large cluster and a number of singletons.

In order to overcome the limitations of the single linkage scheme, in this paper we propose a new linkage criterion called *Genie*. It not only produces high-quality outputs (as compared, e.g., to the Ward and average linkage criteria) but is also relatively fast to compute. The contribution is set out as follows. In the next section, we review some basic properties of hierarchical clustering algorithms and introduce the notion of an inequity index, which can be used to compare some aspects of the quality of clusterings obtained by means of different algorithms. The new linkage criterion, together with its evaluation on diverse benchmark sets, is introduced in Section 3. In Section 4 we propose an algorithm to compute the introduced clustering scheme and test its time performance. Please note that a reference implementation of the Genie method has been included in the `genie` package¹ for R [52]. Finally, Section 5 concludes the paper and provides future research directions.

2. A discussion on classical linkage criteria

While a hierarchical clustering algorithm is being computed on a given data set $\mathcal{X} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}\}$, there are $n - j$ clusters at the j -th step of the procedure, $j = 0, \dots, n - 1$. It is always true that $C^{(j)} = \{C_1^{(j)}, \dots, C_{n-j}^{(j)}\}$ with $C_u^{(j)} \cap C_v^{(j)} = \emptyset$ for $u \neq v$, $C_u^{(j)} \neq \emptyset$, and $\bigcup_{u=1}^{n-j} C_u^{(j)} = \mathcal{X}$. That is, $C^{(j)}$ is a partition of \mathcal{X} .

Initially, we have that $C^{(0)} = \{\{\mathbf{x}^{(1)}\}, \dots, \{\mathbf{x}^{(n)}\}\}$, i.e., $C_i^{(0)} = \{\mathbf{x}^{(i)}\}$, $i = 1, \dots, n$. In other words, each observation is the sole member of its own cluster. When proceeding from step $j - 1$ to j , the clustering procedure decides which of the two clusters $C_u^{(j-1)}$ and $C_v^{(j-1)}$, $u < v$, are to be merged so that we get $C_i^{(j)} = C_i^{(j-1)}$ for $u \neq i < v$, $C_u^{(j)} = C_u^{(j-1)} \cup C_v^{(j-1)}$, and $C_i^{(j)} = C_{i+1}^{(j-1)}$ for $i > v$. In the single (minimum) linkage scheme, u and v are such that:

$$\arg \min_{(u,v), u < v} \left(\min_{\mathbf{a} \in C_u^{(j-1)}, \mathbf{b} \in C_v^{(j-1)}} \mathfrak{d}(\mathbf{a}, \mathbf{b}) \right).$$

On the other hand, the complete (maximum) linkage is based on:

$$\arg \min_{(u,v), u < v} \left(\max_{\mathbf{a} \in C_u^{(j-1)}, \mathbf{b} \in C_v^{(j-1)}} \mathfrak{d}(\mathbf{a}, \mathbf{b}) \right),$$

¹See <https://genieclust.gagolewski.com> for a new implementation – available for both R and Python.

the average linkage on:

$$\arg \min_{(u,v), u < v} \left(\frac{1}{|C_u^{(j-1)}| + |C_v^{(j-1)}|} \sum_{\mathbf{a} \in C_u^{(j-1)}, \mathbf{b} \in C_v^{(j-1)}} \mathfrak{d}(\mathbf{a}, \mathbf{b}) \right),$$

and Ward's (minimum variance) method, compare [41] and also [44], on:

$$\begin{aligned} & \arg \min_{(u,v), u < v} \frac{1}{|C_u^{(j-1)}| + |C_v^{(j-1)}|} \times \\ & \times \left(\sum_{\mathbf{a} \in C_u^{(j-1)}, \mathbf{b} \in C_v^{(j-1)}} 2\mathfrak{d}^2(\mathbf{a}, \mathbf{b}) - \frac{|C_v^{(j-1)}|}{|C_u^{(j-1)}|} \sum_{\mathbf{a}, \mathbf{a}' \in C_u^{(j-1)}} \mathfrak{d}^2(\mathbf{a}, \mathbf{a}') - \frac{|C_u^{(j-1)}|}{|C_v^{(j-1)}|} \sum_{\mathbf{b}, \mathbf{b}' \in C_v^{(j-1)}} \mathfrak{d}^2(\mathbf{b}, \mathbf{b}') \right), \end{aligned}$$

where \mathfrak{d} is a chosen dissimilarity measure.

2.1. Advantages of single-linkage clustering

The main advantage behind the single linkage clustering lies in the fact that its most computationally demanding part deals with solving the minimum spanning tree (MST, see [27]) problem, compare, e.g., the classical Prim's [51] or Kruskal's [37] algorithms as well as a comprehensive historical overview by Graham and Hell [28]. In particular, there is an algorithm [45] which can be run in parallel and which requires exactly $(n^2 - n)/2$ distance computations. Moreover, under certain assumptions on \mathfrak{d} (e.g., the triangle inequality) and the space dimensionality, the Kruskal algorithm may be modified so as to make use of some nearest-neighbor (NN) search data structure which enables to speed up its computations further on (the algorithm can also be run in parallel). Having obtained the MST, a single linkage clustering may then be computed very easily, compare [54].

For the other three mentioned linkage schemes there is, e.g., a quite general nearest-neighbor chains algorithm [43], as well as some other methods which require that, e.g., \mathcal{X} is a subset of the Euclidean space \mathbb{R}^d , see [41, 42] for a survey. Unfortunately, we observe that all these algorithms tend to quite frequently refer to already computed dissimilarities; it may be shown that they use up to $3n^2$ distance computations. Practically, the only way to increase the performance of these algorithms [45] is to pre-compute the whole distance matrix (more precisely, the elements either above or below its diagonal). However, we already noted that such an approach is unusable for n already of moderate order of magnitude ("large data").

2.2. Drawbacks of single-linkage clustering

Nevertheless, it may be observed that unless the underlying clusters are very well-separated, the single linkage approach tends to construct clusters of unbalanced sizes, often




Figure 1: Dendrogram for the single linkage clustering of the *Iris* data set.

resulting – at some fixed dendrogram cut level – in a single large cluster and a number of singletons or ones with a very low cardinality.

For instance, Figure 1 depicts a dendrogram resulting in applying the single linkage clustering on the famous Fisher’s *Iris* data set [19] (available in the R [52] datasets package, object name *iris*) with respect to the Euclidean distance. At the highest level, there are two clusters (50 observations corresponding to *iris setosa* and 100 both to *virginica* and *versicolor*) – these two point groups are well-separated on a 4-dimensional plane. Here, high skewness may be observed in the two subtrees, e.g., cutting the left subtree at the height of 0.4 gives us a partition consisting of three singletons and one large cluster of size 47. These three observations lay slightly further away from the rest of the points. When the $h = 0.4$ cut of the whole tree is considered, there are sixteen singletons, three clusters of size 2, one cluster of size 4, and three large clusters of sizes 38, 39, and 47 (23 clusters in total).

In order to quantitatively capture the mentioned dendrogram *skewness*, we may refer to the definition of an inequity (economic inequality, poverty) index, compare [2, 8, 24] and, e.g., [35, 36] for a different setting.

Definition 1. For a fixed $n \in \mathbb{N}$, let \mathcal{G} denote the set of all non-increasingly ordered n -tuples with elements in the set of non-negative integers, i.e., $\mathcal{G} = \{(x_1, \dots, x_n) \in \mathbb{N}_0^n : x_1 \geq \dots \geq x_n\}$. Then $F : \mathcal{G} \rightarrow [0, 1]$ is an inequity index, whenever:

- (a) it is Schur-convex, i.e., for any $\mathbf{x}, \mathbf{y} \in \mathcal{G}$ such that $\sum_{i=1}^n x_i = \sum_{i=1}^n y_i$, if it holds for all $i = 1, \dots, n$ that $\sum_{j=1}^i x_j \leq \sum_{j=1}^i y_j$, then $F(\mathbf{x}) \leq F(\mathbf{y})$,
- (b) $\inf_{\mathbf{x} \in \mathcal{G}} F(\mathbf{x}) = 0$,
- (c) $\sup_{\mathbf{x} \in \mathcal{G}} F(\mathbf{x}) = 1$.

Notice that, under the current assumptions, if we restrict ourselves to the set of sequences with elements summing up to n , the upper bound (meaning complete inequality) of each inequity index is obtained for $(n, 0, 0, \dots, 0)$ and the lower bound (no inequality) for the $(1, 1, 1, \dots, 1)$ vector.

Every inequity index like F fulfills a crucial property called the Pigou-Dalton principle (also known as *progressive transfers*). Namely, for any $\mathbf{x} \in \mathcal{G}$, $i < j$, and $h > 0$ such that $x_i - h \geq x_{i+1}$ and $x_{j-1} \geq x_j + h$, it holds that:

$$F(x_1, \dots, x_i, \dots, x_j, \dots, x_n) \geq F(x_1, \dots, x_i - h, \dots, x_j + h, \dots, x_n).$$

In other words, any income transfer from a richer to a poorer entity never increases the level of inequity. Notably, such measures of inequality of wealth distribution are not only of interest in economics: it turns out, see [3, 4], that they can be related to ecological indices of evenness [49], which aim to capture how evenly species' populations are distributed over a geographical region, compare [11, 32, 50] or measures of data spread [23].

Among notable examples of inequity measures we may find the normalized Gini-index [25]:

$$G(\mathbf{x}) = \frac{\sum_{i=1}^{n-1} \sum_{j=i+1}^n |x_i - x_j|}{(n-1) \sum_{i=1}^n x_i} \quad (1)$$

or the normalized Bonferroni-index [7]:

$$B(\mathbf{x}) = \frac{n}{n-1} \left(1 - \frac{\sum_{i=1}^n \frac{1}{n-i+1} \sum_{j=i}^n x_j}{\sum_{i=1}^n x_i} \right). \quad (2)$$

Referring back to the above motivational example, we may say that there is often a high inequality between cluster sizes in the case of the single linkage method. Denoting by $c_i = |C_i^{(j)}|$ the size of the i -th cluster at the j -th iteration of the algorithm, $F(c_{(n-j)}, \dots, c_{(1)})$ tends to be very high (here $c_{(i)}$ denotes the i -th smallest value in the given sequence, obviously we always have $\sum_{i=1}^n c_i = n$).

Figure 2 depicts the Gini-indices of the cluster size distributions as a function of the number of clusters in the case of the *Iris* data set and the Euclidean distance. The outputs of four clustering methods are included: single, average, complete, and Ward linkage.

The highest inequality is of course observed in the case of the single linkage algorithm. For instance, if the dendrogram is cut at height of 0.4 (23 clusters in total, their sizes are provided above), the Gini-index is as high as $\simeq 0.76$ (the maximum, 0.85, is obtained for 10 clusters). In this example, the Ward method keeps the Gini-index relatively low. Similar behavior of hierarchical clustering algorithms may be observed for other data sets.




Figure 2: The Gini-indices of the cluster size distributions in the case of the *Iris* data set: single, average, complete, and Ward linkage.

3. The Genie algorithm and its evaluation

3.1. New linkage criterion

In order to compensate the drawbacks of the single linkage scheme, while retaining its simplicity, we propose the following linkage criterion which from now on we refer to as *the Genie algorithm*. Let F be a fixed inequity measure (e.g., the Gini-index) and $g \in (0, 1]$ be some threshold. At step j :

1. if $F(c_{(n-j)}, \dots, c_{(1)}) \leq g$, $c_i = |C_i^{(j)}|$, apply the original single linkage criterion:

$$\arg \min_{(u,v), u < v} \left(\min_{\mathbf{a} \in C_u^{(j)}, \mathbf{b} \in C_v^{(j)}} \mathfrak{d}(\mathbf{a}, \mathbf{b}) \right),$$

2. otherwise, i.e., if $F(c_{(n-j)}, \dots, c_{(1)}) > g$, restrict the search domain only to pairs of clusters such that one of them is of the smallest size:

$$\arg \min_{\substack{(u,v), u < v, \\ |C_u^{(j)}| = \min_i |C_i^{(j)}| \text{ or } \\ |C_v^{(j)}| = \min_i |C_i^{(j)}|}} \left(\min_{\mathbf{a} \in C_u^{(j)}, \mathbf{b} \in C_v^{(j)}} \mathfrak{d}(\mathbf{a}, \mathbf{b}) \right).$$

This modification prevents drastic increases of the chosen inequity measure and forces early merges of small clusters with some other ones. Figure 3 gives the cluster size distribution (compare Figure 2) in case of the proposed algorithm and the *Iris* data set. Here, we used four different thresholds for the Gini-index, namely, 0.3, 0.4, 0.5, and 0.6. Of course, whatever the choice of the inequity index, if $g = 1$, then we obtain the single linkage scheme.




Figure 3: The Gini-indices of the cluster size distributions in the case of the *Iris* data set: the Genie algorithm; the Gini-index thresholds are set to 0.3, 0.4, 0.5, and 0.6.

On a side note, let us point out a small issue that may affect the way the dendrograms resulting in applying the Genie algorithm are plotted. As now the “heights” at which clusters are merged are not being output in a nondecreasing order, they should somehow be adjusted when drawing such diagrams. Yet, the so-called reversals (inversions, departures from ultrametricity) are a well-known phenomenon, see [39], and may also occur in other linkages too (e.g., the nearest-centroid one).

3.2. Benchmark data sets description

In order to evaluate the proposed Genie linkage scheme, we shall test it in various spaces (points in \mathbb{R}^d for some d , images, character strings, etc.), on balanced and unbalanced data of different shapes. Below we describe the 29 benchmark data sets used, 21 of which are in the Euclidean space and the remaining 8 ones are non-Euclidean. All of them are available for download and inspection at <http://www.gagolewski.com/resources/data/clustering/>. Please notice that most of the data sets have already been used in the literature for verifying the performance of various other algorithms.

In each case below, n denotes the number of objects and d – the space dimensionality (if applicable). For every data set its author(s) provided a vector of true (reference) cluster labels. Therefore, we below denote with k the true number of underlying clusters (resulting dendrograms should be cut at this very level during the tests). Moreover, we include the information on whether the reference clusters are of balanced sizes. If this is not the case, the Gini-index of cluster sizes is reported.

Character strings.

- actg1 ($n = 2500$, mean string length $d = 99.9$, $k = 20$, balanced), actg2 ($n = 2500$, mean $d = 199.9$, $k = 5$, the Gini-index of the reference cluster sizes is equal to 0.427), actg3 ($n = 2500$, mean $d = 250.2$, $k = 10$, Gini-index 0.429) – character strings (of varying lengths) over the $\{a, c, t, g\}$ alphabet. First, k random strings (of identical lengths) were generated for the purpose of being cluster centers. Each string in the data set was created by selecting a random cluster center and then performing many Levenshtein edit operations (character insertions, deletions, substitutions) at randomly chosen positions. For use with the Levenshtein distance.
- binstr1 ($n = 2500$, $d = 100$, $k = 25$, balanced), binstr2 ($n = 2500$, $d = 200$, $k = 5$, Gini-index 0.432), binstr3 ($n = 2500$, $d = 250$, $k = 10$, Gini-index 0.379) – character strings (each of the same length d) over the $\{0,1\}$ alphabet. First, k random strings were generated for the purpose of being cluster centers. Each string in the data set was created by selecting a random cluster center and then modifying its digits at randomly chosen positions. For use with the Hamming distance.

Images. These are the first 2000 digits from the famous MNIST database of handwritten digits by Y. LeCun et al., see <http://yann.lecun.com/exdb/mnist/>; clusters are approximately balanced.

- digits2k_pixels ($d = 28 \times 28$, $n = 2000$, $k = 10$) – data consist of 28×28 pixel images. For testing purposes, we use the Hamming distance on corresponding monochrome pixels (color value is marked with 1 if the gray level is in the $(32, 255]$ interval and 0 otherwise).

- `digits2k_points` ($d = 2, n = 2000, k = 10$) – based on the above data set, we represent the contour of each digit as a set of points in \mathbb{R}^2 . Brightness cutoff of 64 was used to generate the data. Each digit was shifted, scaled, and rotated if needed. For testing, we use the Hausdorff (Euclidean-based) distance.

SIPU benchmark data sets. Researchers from the Speech and Image Processing Unit, School of Computing, University of Eastern Finland prepared a list of exemplary benchmarks, which is available at <http://cs.joensuu.fi/sipu/datasets/>. The data sets have already been used in a number of papers. Because of the problems with computing the other linkages in R as well as in Python, see the next section for discussion, we chose only the data sets of sizes ≤ 10000 . Moreover, we omitted the cases in which all the algorithms worked flawlessly, meaning that the underlying clusters were separated too well. In all the cases, we rely on the Euclidean distance.

- `s1` ($n = 5000, d = 2, k = 15$), `s2` ($n = 5000, d = 2, k = 15$), `s3` ($n = 5000, d = 2, k = 15$), `s4` ($n = 5000, d = 2, k = 15$) – S-sets [21]. Reference clusters are more or less balanced.
- `a1` ($n = 3000, d = 2, k = 20$), `a2` ($n = 5250, d = 2, k = 35$), `a3` ($n = 7500, d = 2, k = 50$) – A-sets [38]. Classes are fully balanced.
- `g2-2-100` ($n = 2048, d = 2, k = 2$), `g2-16-100` ($n = 2048, d = 16, k = 2$), `g2-64-100` ($n = 2048, d = 64, k = 2$) – G2-sets. Gaussian clusters of varying dimensions, high variance. Clusters are fully balanced.
- `unbalance` ($n = 6500, d = 2, k = 8$). Unbalanced clusters, the Gini-index of reference cluster sizes is 0.626.
- `Aggregation` ($n = 788, d = 2, k = 7$) [26]. Gini-index 0.454.
- `Compound` ($n = 399, d = 2, k = 6$) [58]. Gini-index 0.440.
- `pathbased` ($n = 300, d = 2, k = 3$) [12]. Clusters are more or less balanced.
- `spiral` ($n = 312, d = 2, k = 3$) [12]. Clusters are more or less balanced.
- `D31` ($n = 3100, d = 2, k = 31$) [55]. Clusters are fully balanced.
- `R15` ($n = 600, d = 2, k = 15$) [55]. Clusters are fully balanced.
- `flame` ($n = 240, d = 2, k = 2$) [22]. Gini-index 0.275.
- `jain` ($n = 373, d = 2, k = 2$) [33]. Gini-index 0.480.

Iris. The Fisher’s *Iris* [19] data set, available in the R [52] `datasets` package. Again, the Euclidean distance is used.

- `iris` ($d = 4, n = 150, k = 3$) – the original data set. Fully balanced clusters.
- `iris5` ($d = 4, n = 105, k = 3$) – an unbalanced version of the above one, in which we took only five last observations from the first group (*iris setosa*). Gini-index 0.429.

3.3. Benchmark results

In order to quantify the degree of agreement between two k -partitions of a given set, the notion of the FM-index [20] is very often used.

Definition 2. Let $C = \{C_1, \dots, C_k\}$ and $C' = \{C'_1, \dots, C'_k\}$ be two k -partitions of the set $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}\}$. The Fowlkes-Mallows (FM) index is given by:

$$FM\text{-index}(C, C') = \frac{\sum_{i=1}^k \sum_{j=1}^k m_{i,j}^2 - n}{\sqrt{\left(\sum_{i=1}^k \left(\sum_{j=1}^k m_{i,j}\right)^2 - n\right) \left(\sum_{j=1}^k \left(\sum_{i=1}^k m_{i,j}\right)^2 - n\right)}} \in [0, 1],$$

where $m_{i,j} = |C_i \cap C'_j|$.

If the two partitions are equivalent (equal up to a permutation of subsets in one of the k -partitions), then the FM-index is equal to 1. Moreover, if each pair of objects that appear in the same set in C appear in two different sets in C' , then the index is equal to 0.

Let us compare the performance of the Genie algorithm (with F set to be the Gini-index; five different thresholds, $g \in \{0.2, 0.3, \dots, 0.6\}$, are used) as well as the single, average, complete, and Ward linkage schemes (as implemented in the `hclust()` function from the R [52] package `stats`). In order to do so, we compute the values of $FM\text{-index}(C, C')$, where C denotes the vector of true (reference) cluster labels (as described in Section 3.2), while C' is the clustering obtained by cutting at an appropriate level the dendrogram returned by a hierarchical clustering algorithm being investigated. However, please observe that for some benchmark data sets the distance matrices consist of non-unique elements. As a result, the output of the algorithms may vary slightly from call to call (this is the case of all the tested methods). Therefore, we shall report the median FM-index across 10 runs of randomly permuted observations in each benchmark set.

Table 1 gives the FM-indices for the 9 clustering methods and the 29 benchmark sets. Best results are marked with bold font. Aggregated basic summary statistics (minimum, quartiles, maximum, arithmetic mean, and standard deviation) for all the benchmark sets are provided in Table 2. Moreover, Figure 4 depicts violin plots of the FM-index distribution².

The highest mean and median FM scores were obtained for the Genie algorithm with a threshold of $g = 0.2$. This setting also leads to the best minimal (worst-case) FM-index. A general observation is that all the tested Gini-index thresholds gave the lowest variance in the FM-indices.

²A violin depicts a box-and-whisker plot (boxes range from the 1st to the 3rd quartile, the median is marked with a white dot) together with a kernel density estimator of the empirical distribution.

Table 1: FM-indices for the 29 benchmark sets and the 9 hierarchical clustering methods studied.

benchmark	single	complete	Ward	average	gini_0.2	gini_0.3	gini_0.4	gini_0.5	gini_0.6
actg1	0.222	0.997	0.998	0.998	0.996	0.941	0.947	0.905	0.624
actg2	0.525	1.000	1.000	1.000	0.975	0.975	0.976	1.000	1.000
actg3	0.383	1.000	1.000	1.000	0.884	0.975	0.975	1.000	0.840
binstr1	0.198	0.874	0.942	0.947	0.952	0.908	0.863	0.749	0.542
binstr2	0.525	0.989	0.994	0.994	0.907	0.909	0.965	0.965	0.819
binstr3	0.368	0.946	0.969	0.971	0.832	0.931	0.937	0.811	0.692
digits2k_pixels	0.315	0.310	0.561	0.326	0.584	0.473	0.473	0.368	0.321
digits2k_points	0.315	0.256	0.458	0.280	0.671	0.601	0.559	0.438	0.405
s1	0.589	0.973	0.984	0.983	0.989	0.989	0.989	0.989	0.989
s2	0.257	0.807	0.912	0.918	0.921	0.921	0.791	0.804	0.767
s3	0.257	0.548	0.699	0.636	0.708	0.690	0.610	0.609	0.559
s4	0.257	0.468	0.585	0.546	0.644	0.620	0.563	0.529	0.482
a1	0.564	0.920	0.918	0.929	0.940	0.905	0.901	0.849	0.776
a2	0.480	0.911	0.924	0.936	0.951	0.925	0.903	0.843	0.703
a3	0.449	0.919	0.939	0.945	0.958	0.940	0.923	0.836	0.743
g2-2-100	0.707	0.586	0.598	0.706	0.601	0.602	0.637	0.648	0.648
g2-16-100	0.707	0.897	0.923	0.707	0.842	0.697	0.697	0.697	0.704
g2-64-100	0.707	1.000	1.000	0.707	0.999	0.999	0.999	0.999	0.999
unbalance	0.999	0.775	1.000	1.000	0.723	0.730	0.775	0.844	0.911
Aggregation	0.861	0.833	0.842	1.000	0.582	0.657	0.816	0.908	0.894
Compound	0.830	0.855	0.653	0.862	0.638	0.649	0.637	0.708	0.889
pathbased	0.573	0.595	0.674	0.653	0.751	0.751	0.751	0.751	0.751
spiral	1.000	0.339	0.337	0.357	1.000	1.000	1.000	1.000	1.000
D31	0.349	0.926	0.923	0.910	0.937	0.903	0.828	0.742	0.695
R15	0.637	0.980	0.983	0.990	0.987	0.987	0.987	0.823	0.637
flame	0.730	0.623	0.624	0.731	1.000	1.000	1.000	1.000	1.000
jain	0.804	0.922	0.790	0.922	1.000	1.000	1.000	1.000	1.000
iris	0.764	0.769	0.822	0.841	0.923	0.923	0.923	0.923	0.754
iris5	0.691	0.665	0.738	0.765	0.764	0.764	0.764	0.886	0.673
FM-index ≥ 0.95	2	7	9	9	11	8	9	8	6
FM-index ≥ 0.9	2	13	16	16	16	18	15	11	7
FM-index < 0.7	19	9	9	6	6	8	7	6	11
FM-index < 0.5	12	4	2	3	0	1	1	2	3

Table 2: Basic summary statistics of the FM-index distribution over the 29 benchmark sets.

	single	complete	Ward	average	gini_0.2	gini_0.3	gini_0.4	gini_0.5	gini_0.6
Min	0.198	0.256	0.337	0.280	0.582	0.473	0.473	0.368	0.321
Q1	0.349	0.623	0.674	0.707	0.723	0.697	0.751	0.742	0.648
Median	0.564	0.874	0.918	0.918	0.921	0.909	0.901	0.843	0.751
Q3	0.707	0.946	0.983	0.983	0.975	0.975	0.975	0.965	0.894
Max	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000
Mean	0.554	0.782	0.820	0.812	0.850	0.840	0.834	0.815	0.752
St.Dev.	0.235	0.225	0.188	0.215	0.147	0.156	0.159	0.172	0.185




Figure 4: Violin plots of the FM-index distribution over the 29 benchmark sets.

Table 3: Basic summary statistics of the FM-index distribution over the 21 Euclidean benchmark sets.

	single	complete	ward	average	gini_0.2	gini_0.3	gini_0.4	gini_0.5	gini_0.6	BIRCH	k-means
Min	0.257	0.339	0.337	0.357	0.582	0.602	0.563	0.529	0.482	0.350	0.327
Q1	0.480	0.623	0.674	0.707	0.723	0.697	0.751	0.742	0.695	0.653	0.701
Median	0.691	0.833	0.842	0.862	0.923	0.905	0.828	0.843	0.754	0.894	0.821
Q3	0.764	0.920	0.924	0.936	0.987	0.987	0.987	0.923	0.911	0.924	0.969
Max	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000
Mean	0.629	0.777	0.803	0.812	0.850	0.841	0.833	0.828	0.789	0.801	0.816
St.Dev.	0.224	0.187	0.177	0.172	0.150	0.146	0.145	0.138	0.156	0.183	0.177

It is of course unsurprising that there is no free lunch in data clustering – no algorithm is perfect on all the data sets. All the tested hierarchical clustering algorithms were far from perfect ($FM < 0.7$) on the `digits2k_pixels`, `digits2k_points`, and `s4` data sets. However, in overall, the single linkage clustering is particularly bad (except for the `unbalance` and `spiral` data sets). Among the other algorithms, the complete linkage and the Genie algorithm for $g \geq 0.5$ give the lowest average and median FM-index. All the other methods (Genie with thresholds of $g < 0.5$, Ward, average linkage) are very competitive. Also please keep in mind that for the Genie algorithm with a low inequity index threshold we expect a loss in performance for unbalanced clusters sizes

Finally, let us compare the performance of the 9 hierarchical clustering algorithms as well as the k -means and BIRCH (`threshold=0.5`, `branching_factor=10`) procedures (both implemented in the `scikit-learn` package for Python). Now we are of course limited only to data in the Euclidean space, therefore the number of benchmark data sets reduces to 21. Table 3 gives basic summary statistics of the FM-index distributions. We see that in this case the Genie algorithm ($g < 0.5$) outperforms all the methods being compared too.

Taking into account our algorithm’s out-standing performance and – as it shall turn out in the next section – relatively low run-times (especially on larger data sets and compared with the average or Ward linkage), the proposed method may be recommended for practical use.

4. Possible implementations of the Genie linkage algorithm

Having shown the high usability of the new approach, let us discuss some ways to implement the Genie clustering method in very detail. We have already stated that the most important part of computing the single linkage algorithm consists of determining a minimal spanning tree (this can be non-unique if there are pairs of objects with identical dissimilarity degrees) of the complete undirected weighted graph corresponding to objects in \mathcal{X} and the pairwise dissimilarities. It turns out that we have what follows.

Theorem 1. *The Genie linkage criterion can be implemented based on an MST.*

Sketch of the proof. By [51, Principle 2], in order to construct a minimal spanning tree, it is sufficient to connect any two disconnected minimal spanning subtrees via an edge of minimal weight and iterate such a process until a single connected tree is obtained. As our linkage criterion (Section 3.1) always chooses such an edge, the proof is complete. \square

```

0. Input:  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$  –  $n$  objects,  $g \in (0, 1]$  – inequity index threshold,
    $\mathfrak{d}$  – a dissimilarity measure;
1.  $ds = \text{DisjointSets}(\{1\}, \{2\}, \dots, \{n\})$ ;
2.  $m = \text{MST}(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)})$ ; /* see Theorem 1 */
3.  $pq = \text{MinPriorityQueue}<\text{PQItem}>(\emptyset)$ ;
   /* PQItem structure: (index1, index2, dist);
   pq returns the element with the smallest dist */
4. for each weighted edge  $(i, j, d_{i,j})$  in  $m$ :
   4.1.  $pq.\text{push}(\text{PQItem}(i, j, d_{i,j}))$ ;
5. for  $j = 1, 2, \dots, n - 1$ :
   5.1. if  $ds.\text{compute\_inequity}() \leq g$ : /* e.g., the Gini-index */
     5.1.1.  $t = pq.\text{pop}()$ ; /* PQItem with the least dist */
     else:
       5.1.2.  $t = pq.\text{pop\_conditional} \left( t: ds.\text{size}(t.\text{index1}) = ds.\text{min\_size}() \right.$ 
or  $ds.\text{size}(t.\text{index2}) = ds.\text{min\_size}()$ );
/* PQItem with the least dist that fulfills the given logical condition */
   5.2.  $s_1 = ds.\text{find\_set}(t.\text{index1})$ ;
   5.3.  $s_2 = ds.\text{find\_set}(t.\text{index2})$ ; /* assert:  $s_1 \neq s_2$  */
   5.4. output “linking  $(s_1, s_2)$ ”;
   5.5.  $ds.\text{link}(t.\text{index1}, t.\text{index2})$ ;

```

Figure 5: A pseudocode for the Genie algorithm.

Please note that, by definition, the weighted edges appear in an MST in no particular order – for instance, the Prim [51] algorithm’s output depends on the permutation of inputs. Therefore, having established the above relation between the Genie clustering and an MST, in Figure 5 we provide a pseudocode of the algorithm that guarantees the right cluster

merge order. The procedure resembles the Kruskal [37] algorithm and is fully concordant with our method’s description in Section 3.1.

Observe that the very same algorithm can be used to compute the single linkage clustering (in such a case, step 5.1 – in which we compute a chosen inequity measure – as well as step 5.1.2 – in which we test whether the inequity measure raises above a given threshold – are never executed).

Example 1. *Let us consider an exemplary data set consisting of 5 points in the real line: $\mathbf{x}^{(1)} = 0, \mathbf{x}^{(2)} = 1, \mathbf{x}^{(3)} = 3, \mathbf{x}^{(4)} = 6, \mathbf{x}^{(5)} = 10$ and let us set \mathfrak{d} to be the Euclidean distance. The minimum spanning tree – which, according to Theorem 1 gives us complete information needed to compute the resulting clustering – consists of the following edges: $\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}\}, \{\mathbf{x}^{(2)}, \mathbf{x}^{(3)}\}, \{\mathbf{x}^{(3)}, \mathbf{x}^{(4)}\}, \{\mathbf{x}^{(4)}, \mathbf{x}^{(5)}\}$, with weights 1, 2, 3, and 4, respectively.*

Let the inequity measure F be the Gini-index. If the threshold g is set to 1.0 (the single linkage), the merge steps are as follows. Please note that the order in which we merge the clusters is simply determined by sorting the edges in the MST increasingly by weights, just as in the Kruskal algorithm.

step	current partitioning	MST edge	Gini-index
0.	$\{\mathbf{x}^{(1)}\}, \{\mathbf{x}^{(2)}\}, \{\mathbf{x}^{(3)}\}, \{\mathbf{x}^{(4)}\}, \{\mathbf{x}^{(5)}\}$	$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}\}_1$	0.0
1.	$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}\}, \{\mathbf{x}^{(3)}\}, \{\mathbf{x}^{(4)}\}, \{\mathbf{x}^{(5)}\}$	$\{\mathbf{x}^{(2)}, \mathbf{x}^{(3)}\}_2$	0.2
2.	$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \mathbf{x}^{(3)}\}, \{\mathbf{x}^{(4)}\}, \{\mathbf{x}^{(5)}\}$	$\{\mathbf{x}^{(3)}, \mathbf{x}^{(4)}\}_3$	0.4
3.	$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \mathbf{x}^{(3)}, \mathbf{x}^{(4)}\}, \{\mathbf{x}^{(5)}\}$	$\{\mathbf{x}^{(4)}, \mathbf{x}^{(5)}\}_4$	0.6
4.	$\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \mathbf{x}^{(3)}, \mathbf{x}^{(4)}, \mathbf{x}^{(5)}\}$	—	—

However, if the threshold is set to $g = 0.3$, then when proceeding from step 2 to step 3, we need to link $\{\mathbf{x}^{(4)}\}$ and $\{\mathbf{x}^{(5)}\}$ instead of $\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \mathbf{x}^{(3)}\}$ and $\{\mathbf{x}^{(4)}\}$ – that is, the merge is based now on a different MST edge: $\{\mathbf{x}^{(4)}, \mathbf{x}^{(5)}\}_4$ instead of $\{\mathbf{x}^{(3)}, \mathbf{x}^{(4)}\}_3$. Therefore, the resulting 2-partition will be different than the above one: we obtain $\{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \mathbf{x}^{(3)}\}, \{\mathbf{x}^{(4)}, \mathbf{x}^{(5)}\}$ (Gini-index = 0.2).

4.1. Implementation details

In this study, we decided to focus on the Gini-index. In order to make the algorithm time-efficient, this inequity index must be computed incrementally. Let g_j denote the Gini-index value at the time when there are $n - j$ clusters, $j = 0, 1, \dots, n - 1$. Initially, when each cluster is of size 1, the Gini-index is of course equal to 0, hence $g_0 = 0$. Assume that when proceeding from step $j - 1$ to j we link clusters of sizes c_{s_1} and c_{s_2} . It can easily be shown that:

$$g_j = \frac{(n - j) n g_{j-1} + \sum_{i=1}^{n-j+1} (|c_i - c_{s_1} - c_{s_2}| - |c_i - c_{s_1}| - |c_i - c_{s_2}|) - c_{s_2} - c_{s_1} + |c_{s_1} - c_{s_2}|}{(n - j - 1) n}.$$

In other words, after each merge operation, the index is updated, which requires $O(n)$ operations instead of $O(n^2)$ when the index is recomputed from scratch based on the original formula (1). On a side note, it might be shown that the Gini-index can be written as a linear combination of order statistics (compare, e.g., the notion of an OWA operator), but the use of such a definition would require sorting the cluster size vector in each iteration or relying on an ordered search tree-like data structure.

Moreover, please note that the algorithm is based on a non-standard implementation of the disjoint sets data structure. Basically, the required extension that keeps track of cluster counts can be implemented quite easily.

The implementation of the `pop_conditional()` method in the priority queue *pq* (we used a heap-based data structure) can be done by means of an auxiliary queue, onto which elements not fulfilling the logical condition given in step 5.1.2 are temporarily moved. If, after a merge operation, the inequity index is still above the desired threshold and the minimal cluster size did not change since the previous iteration, the auxiliary priority queue is kept as-is and we continue to seek a cluster of the lowest cardinality from the same place. Otherwise, the elements must go back to *pq* and the search must start over.

To conclude, the cost of applying the cluster merge procedure ($O(n^2)$ pessimistically for $g < 1$ and $O(n \log n)$ for $g = 1$) is not dominated by the cost of determining an MST ($O(Cn^2)$ pessimistically, under our assumptions this involves a function C of data dimensionality and reflects the cost of computing a given dissimilarity measure \mathfrak{d}). Hence, the new clustering scheme shall give us run-times that are comparable with the single linkage method. It is worth noting that the time complexity ($\Theta(n^2)$) as well as the memory complexity ($\Theta(n)$) of the algorithm is optimal (as far as the whole class of hierarchical clustering algorithms is concerned, compare [41]).

On a side note, the NN-chains algorithm, which is suitable for solving – among others – the complete, average, and Ward linkage clustering also has a time complexity of $O(n^2)$ [41, 43], but, as we shall see in further on, it requires computing at least 2–5 times more pairwise distances.

For the sake of comparison, let us study two different algorithms to determine an MST.

1. An MST algorithm based on Prim's one. A first algorithm, sketched in [45, Figure 8], is quite similar to the one by Prim [51]. Its pseudocode is given in Figure 6. Note that the algorithm guarantees that exactly $(n^2 - n)/2$ pairwise distances are computed. Moreover, the inner loop can be run in parallel. In such a case, M should be a vector of indices not yet in the MST – due to that the threads may have random access to such an array – and step 6.2.3 should be moved to a separate loop – so as to a costly critical section is avoided.

```

0. Input:  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$  –  $n$  objects,  $\mathfrak{d}$  – a dissimilarity measure;
1.  $F = (\infty, \dots, \infty)$  ( $n$  times);           /*  $F_j$  – index of  $j$ 's nearest neighbor */
2.  $D = (\infty, \dots, \infty)$  ( $n$  times);       /*  $D_j$  – distance to  $j$ 's nearest neighbor */
3.  $lastj = 1$ ;
4.  $m = \emptyset$ ;                               /* a resulting MST */
5.  $M = \{2, 3, \dots, n\}$ ;                     /* indices not yet in  $m$  */
6. for  $i = 1, 2, \dots, n - 1$ :
    6.1.  $bestj = 1$ ;
    6.2. for each  $j \in M$ :
        6.2.1.  $d = \mathfrak{d}(\mathbf{x}^{(lastj)}, \mathbf{x}^{(j)})$ ;
        6.2.2. if  $d < D_j$ :
            6.2.2.1.  $D_j = d$ ;
            6.2.2.2.  $F_j = lastj$ ;
        6.2.3. if  $D_j < D_{bestj}$ :                 /* assert:  $D_1 = \infty$  */
            6.2.3.1.  $bestj = j$ ;
    6.3.  $m = m \cup \{(F_{bestj}, bestj, D_{bestj})\}$ ; /* add an edge to  $m$  */
    6.4.  $M = M \setminus \{bestj\}$ ;               /* now this index is in  $m$  */
    6.5.  $lastj = bestj$ ;
7. return  $m$ ;

```

Figure 6: A simple $(n^2 - n)/2$ algorithm to determine an MST.

2. *An MST algorithm based on Kruskal's one.* The second algorithm considered is based on the one by Kruskal [37] and its pseudocode is given in Figure 7. It relies on a method called `getNextNearestNeighbor()`, which fetches the index j of the next not-yet considered nearest (in terms of increasing \mathfrak{d}) neighbor of an object at index i having the property that $j > i$. If such a neighbor does not exist anymore, the function returns ∞ . Please observe that in the prefetch phase the calls to `getNextNearestNeighbor()` can be run in parallel.

A naïve implementation of the `getNextNearestNeighbor()` function requires either $O(1)$ time and $O(n)$ memory for each object (i.e., $O(n^2)$ in total – the lists of all neighbors can be stored in n priority queues, one per each object) or $O(n)$ time and $O(1)$ memory (i.e., $O(n)$ in total – no caching done at all). As our priority is to retain total $O(n)$ memory

use, the mentioned approach is of course expected to have a much worse time performance than the Prim-based one.

However, now let us assume that a dissimilarity measure $\mathfrak{d} : \mathcal{X} \times \mathcal{X} \rightarrow [0, \infty]$ is in fact a pseudometric, i.e., it additionally fulfills the triangle inequality: for any $\mathbf{x}, \mathbf{y}, \mathbf{z} \in \mathcal{X}$ we have $\mathfrak{d}(\mathbf{x}, \mathbf{y}) \leq \mathfrak{d}(\mathbf{x}, \mathbf{z}) + \mathfrak{d}(\mathbf{z}, \mathbf{y})$ – such a setting often occurs in practice.

```

0. Input:  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}$  –  $n$  objects,  $\mathfrak{d}$  – a dissimilarity measure;
1.  $pq = \text{MinPriorityQueue}<\text{PQItem}>(\emptyset);$ 
   /* PQItem structure: (index1, index2, dist);
   pq returns the element with the smallest dist */
2. for  $i = 1, 2, \dots, n - 1$ : /* prefetch phase */
   2.1.  $j = \text{getNextNearestNeighbor}(i);$  /* depends on  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}, \mathfrak{d}$  */
   2.2.  $pq.\text{push}(\text{PQItem}(i, j, \mathfrak{d}(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})));$  /* assert:  $i < j$  */
3.  $ds = \text{DisjointSets}(\{1\}, \{2\}, \dots, \{n\});$ 
4.  $m = \emptyset;$ 
5.  $i = 1;$ 
6. while  $i < n$ : /* merge phase */
   6.1.  $t = pq.\text{pop}();$  /* PQItem with the least dist */
   6.2.  $s_1 = ds.\text{find\_set}(t.\text{index1});$ 
   6.3.  $s_2 = ds.\text{find\_set}(t.\text{index2});$ 
   6.4. if  $s_1 \neq s_2$ :
     6.4.1.  $i = i + 1;$ 
     6.4.2.  $m = m \cup \{(t.\text{index1}, t.\text{index2}, t.\text{dist})\};$ 
     6.4.3.  $ds.\text{link}(t.\text{index1}, t.\text{index2});$ 
   6.5.  $j = \text{getNextNearestNeighbor}(t.\text{index1});$  /* assert:  $t.\text{index1} < j$  */
   6.6. if  $j < \infty$ :
     6.6.1.  $pq.\text{push}(\text{PQItem}(t.\text{index1}, j, \mathfrak{d}(\mathbf{x}^{(t.\text{index1})}, \mathbf{x}^{(j)})));$ 
7. return  $m;$ 

```

Figure 7: A getNextNearestNeighbor()-based algorithm to determine an MST.

In such a case, a significant speed up may be obtained by relying on some nearest-neighbor search data structure supporting queries like “fetch a few nearest-neighbors of the i -th object within the distance in range $[r_{\min}, r_{\max})$ ”, for some $r_{\min} < r_{\max}$. Of course, it is widely known that – due to the so-called curse of dimensionality, compare [1, 5, 9, 13, 53]

– there is no general-purpose algorithm which always works better than the naïve method in spaces of high dimension. Nevertheless, in our case, some modifications of a chosen data structure may lead to improvements in time performance.

Our carefully tuned-up reference implementation (discussed below) is based on a vantage point (VP)-tree, see [57]. The most important modifications applied are as follows.

- Each tree node stores an information on the maximal object index that can be found in its subtrees. This speeds up the search for NNs of objects with higher indices. No distance computation are performed for a pair of indices (i, j) unless $i < j$.
- Each tree node includes an information whether all its subtrees store elements from the same set in the disjoint sets data structure ds . This Boolean flag is recursively updated during a call to `getNextNearestNeighbor()`. Due to that, a significant number of tree nodes during the merge phase can be pruned.
- An actual tree query returns a batch of nearest-neighbors of adaptive size between 20 and 256 (the actual count is determined automatically according to how the underlying VP-tree prunes the child nodes during the search). The returned set of nearest-neighbors is cached in a separate priority queue, one per each input data point. Note that the size of the returned batch guarantees asymptotic linear total memory use.

4.2. The *genie* package for R

A reference implementation of the Genie algorithm has been included in the *genie* package for R [52]. This software is distributed under the open source GNU General Public License, version 3. The package is available for download at the official CRAN (Comprehensive R Archive Network) repository, see <https://cran.r-project.org/web/packages/genie/>, and hence can be installed from within an R session via a call to `install.packages("genie")`. All the core algorithms have been developed in the C++11 programming language; the R interface has been provided by means of the *Rcpp* [17] package. What is more, we decided to rely on the OpenMP API in order to enable multi-threaded computations.

A data set's clustering can be determined via a call to the `genie::hclust2()` function. The `objects` argument, with which we provide a data set to be clustered, may be a numeric matrix, a list of integer vectors, or an R character vector. The dissimilarity measure is selected via the `metric` argument, e.g., "euclidean", "manhattan", "maximum", "hamming", "levenshtein", "dinu", etc. The `thresholdGini` argument can be used to define the threshold for the Gini-index (denoted with g in Figure 5). Finally, the `useVpTree` argument can be used to switch between the MST algorithms given

in Figures 6 (the default) and 7. For more details, please refer to the function’s manual page (`?genie::hclust2`).

Here is an exemplary R session in which we compute the clustering of the flame data set.

```
# load the 'flame' benchmark data set,
# see http://www.gagolewski.com/resources/data/clustering/
data <- as.matrix(read.table(gzfile("flame.data.gz")))
labels <- scan(gzfile("flame.labels.gz"), quiet=TRUE)

# run the Genie algorithm, threshold g=0.2
result <- genie::hclust2(objects=data, metric="euclidean",
  thresholdGini=0.2)

# get the number of reference clusters
k <- length(unique(labels))

# plot the results
plot(data[,1], data[,2], col=labels, pch=cutree(result, k))

# compute the FM-index
as.numeric(dendextend::FM_index(labels, cutree(result, k),
  include_EV=FALSE))
## [1] 1
```

4.3. Number of calls to the dissimilarity measure

Let us compare the number of calls to the dissimilarity measure \mathfrak{d} required by different clustering algorithms. The measures shall be provided relative to $(n^2 - n)/2$, which is denoted with “100%”.

The benchmark data sets are generated as follows. For a given n , σ , and d , $k = 10$ cluster centers $\boldsymbol{\mu}^{(1)}, \dots, \boldsymbol{\mu}^{(k)}$ are picked randomly from the uniform distribution on $[0, 10]^d$. Then, each of the n observations is generated as $\boldsymbol{\mu}^{(j)} + \mathbf{y}$, where y_l for each $l = 1, \dots, d$ is a random variate from the normal distribution with expectation of 0 and standard deviation of σ and j is a random number distributed uniformly in $\{1, 2, \dots, k\}$. In other words, such a data generation method is more or less equivalent to the one used in case of the g2 data sets in the previous section. Here, \mathfrak{d} is set to be the Euclidean metric.

Table 4: Relative number of pairwise distance computations (as compared to $(n^2 - n)/2$) together with FM-indices (in parentheses).

σ	d	n	gini_0.3	gini_1.0 (single)	complete	Ward	average
0.50	2	10000	*4.8% (0.76)	100% (0.38)	476% (0.72)	204% (0.80)	484% (0.78)
0.50	5	10000	*22.0% (1.00)	100% (0.86)	493% (1.00)	221% (1.00)	496% (1.00)
1.50	10	10000	*30.3% (0.96)	100% (0.32)	496% (0.98)	240% (0.98)	499% (0.91)
1.50	15	10000	*58.3% (1.00)	100% (0.42)	497% (1.00)	253% (1.00)	498% (0.98)
1.50	20	10000	*84.9% (1.00)	100% (0.69)	497% (1.00)	261% (1.00)	498% (1.00)
3.50	100	10000	*101.8% (1.00)	100% (0.32)	498% (1.00)	299% (1.00)	499% (1.00)
5.00	250	10000	*100.9% (0.99)	100% (0.32)	498% (1.00)	312% (1.00)	499% (1.00)
1.5	10	100000	*14.1% (0.94)	100% (0.32)	—	241% (0.98)	—
3.5	100	100000	*104.0% (0.98)	100% (0.32)	—	† 321% (0.99)	—

(*) – a VP-tree used (Fig. 7) to determine an MST; the $(n^2 - n)/2$ algorithm (Fig. 6) could always be used instead.

(†) – only one run of the experiment was conducted.

In each test case (different choices of n , d , σ), we randomly generated ten different data sets and averaged the resulting FM-indices and relative numbers of calls to `b`. Table 4 compares:

- the Genie algorithm (`genie::hclust2()`, package version 1.0.0) based on each of the aforementioned MST algorithms (`useVpTree` equals either to `TRUE` or `FALSE`); the Gini-index threshold of 0.3 and 1.0, the latter is equivalent to the single linkage criterion,
- the Ward linkage (the `hclust.vector()` function from the `fastcluster` 1.1.16 package [42] – this implementation works only in the case of the Euclidean distance but uses $O(n)$ memory),
- the complete and average linkage (`fastcluster::hclust()` – they use an NN-chains-based algorithm, and require a pre-computed distance matrix, therefore utilizing $O(n^2)$ memory).

We observe a positive impact of using a metric tree data structure (`useVpTree=TRUE`) in low-dimensional spaces. In high-dimensional spaces, it is better to rely on the $(n^2 - n)/2$ (Prim-like) algorithm. Nevertheless, we observe that in high dimensional spaces the relative number of calls to the dissimilarity measure is 2–5 times smaller than in the case of other linkages. Please note that an NN-chains-based version of the Ward linkage (not listed

in the table; `fastcluster::hclust()` gives similar results as the complete and average ones and that its Euclidean-distance specific version (`fastcluster::hclust.vector()`) seems to depend on the data set dimensionality.

4.4. Exemplary run-times

Let us inspect exemplary run-time measurements of the `genie::hclust2()` function (genie package version 1.0.0). The measurements were performed on a laptop with a Quad-core Intel(R) Core(TM) i7-4700HQ @ 2.40GHz CPU and 16 GB RAM. The computer was running Fedora 21 Linux (kernel 4.1.13-100) and the gcc 4.9.2 compiler was used (`-O2 -march=native` optimization flags).

Table 5: Exemplary run-times (in seconds) for different thread numbers, $n = 100,000$.

data	MST algorithm	g	Number of threads		
			1	2	4
$d = 10, \sigma = 1.5$	Fig. 7 with a VP-tree	0.3	46.5	33.4	28.2
	Fig. 6	0.3	91.5	59.9	44.8
	Fig. 7 with a VP-tree	1.0	32.0	19.1	13.4
	Fig. 6	1.0	77.5	47.7	31.6
$d = 100, \sigma = 3.5$	Fig. 7 with a VP-tree	0.3	1396	740	456
	Fig. 6	0.3	743	413	293
	Fig. 7 with a VP-tree	1.0	1385	717	454
	Fig. 6	1.0	734	396	281

Table 5 summarizes the results for $n = 100,000$ and 1, 2, as well as 4 threads (set up via the `OMP_THREAD_LIMIT` environmental variable). The experimental data sets were generated in the same manner as above. The reported times are minimums of 3 runs. Note that the results also include the time needed to generate some additional objects, so that the output is of the same form as the one generated by the `stats::hclust()` function in R.

We note that running 4 threads at a time (on a single multi-core CPU) gives us a 2–3-fold speed-up. Moreover, a VP-tree-based implementation (Figure 7, `useVpTree=TRUE`) is twice as costly as the other one in spaces of high dimensions. However, in spaces of low dimension it outperforms the $(n^2 - n)/2$ approach (Figure 6). Nevertheless, if a user is unsure whether he/she deals with a high- or low-dimensional space and n is of moderate order of magnitude, the simple approach should rather be recommended, as it gives much more predictable timings. This is why we have decided that the `useVpTree` argument should default to `FALSE`.

For a point of reference, let us note that a single test run of the Ward algorithm (`fastcluster::hclust.vector()`, single-threaded) for $n = 100,000$, $d = 10$, $\sigma = 1.5$ required 1452.8 seconds and for $d = 100$, $\sigma = 3.5$ – as much as 18433.7 seconds (which is almost 25 times slower than the Genie approach).

5. Conclusions

We have presented a new hierarchical clustering linkage criterion which is based on the notion of an inequity (poverty) index. The performed benchmarks indicate that the proposed algorithm – unless the underlying cluster structure is drastically unbalanced – works in overall better not only than the widely used average and Ward linkage scheme, but also than the k -means and BIRCH algorithms which can be applied on data in the Euclidean space only.

Our method requires up to $(n^2 - n)/2$ distance computations, which is ca. 2–5 times less than in the case of the other popular linkage schemes. Its performance is comparable with the single-linkage clustering. As there is no need to store the full distance matrix, the algorithm can be used to cluster larger (within one order of magnitude) data sets than with the Ward and average linkage schemes.

Nevertheless, it seems that we have reached a kind of general limit of an input data set size for “classical” hierarchical clustering, especially in case of multidimensional data. Due to the curse of dimensionality, we do not have any nearest-neighbor search data structures that would enable us to cluster data sets of sizes greater than few millions of observations in a reasonable time span. What is more, we should keep in mind that the lower bound for run-times of all the hierarchical clustering methods is $\Omega(n^2)$ anyway. However, let us stress that for smaller (non-big-data) samples, hierarchical clustering algorithms are still very useful. This is due to the fact that they do not require a user to provide the desired number of clusters in advance and that only a measure of objects’ dissimilarity – fulfilling very mild properties – must be provided in order to determine a data partition.

Further research on the algorithm shall take into account the effects of, e.g., choosing different inequity measures or relying on approximate nearest-neighbors search algorithms and data dimension reduction techniques on the clustering quality. Moreover, in a distributed environment, one may consider partitioning subsets of input data individually and then rely on some clustering aggregation techniques, compare, e.g., [26].

Finally, let us note that the Genie algorithm depends on a free parameter, namely, the inequity index merge threshold, g . The existence of such a tuning parameter is an advantage, as a user may select its value to suit her/his needs. In the case of the Gini-index, we recommend the use of $g \in [0.2, 0.5)$, depending on our knowledge of the underlying cluster distribution. Such a choice led to outstanding results during benchmark studies.

However, we should keep in mind that if the threshold is too low, the algorithm might have problems with correctly identifying clusters of smaller sizes in case of unbalanced data. On the other hand, g cannot be too large, as the algorithm might start to behave as the single-linkage method, which has a very poor performance. A possible way to automate the choice of g could consist of a few pre-flight runs (for different thresholds) on a randomly chosen data sample, a verification of the obtained preliminary clusterings' qualities, and a choice of the best coefficient for the final computation.

Acknowledgments

We would like to thank the Anonymous Reviewers for the constructive comments that helped to significantly improve the manuscript's quality. Moreover, we are indebted to Łukasz Błaszczuk for providing us with the `scikit-learn` algorithms performance results.

This study was supported by the National Science Center, Poland, research project 2014/13/D/HS4/01700. Anna Cena and Maciej Bartoszek would like to acknowledge the support by the European Union from resources of the European Social Fund, Project PO KL "Information technologies: Research and their interdisciplinary applications", agreement UDA-POKL.04.01.01-00-051/10-00 via the Interdisciplinary PhD Studies Program.

References

- [1] C.C. Aggarwal, A. Hinneburg, D.A. Keim, On the surprising behavior of distance metric in high-dimensional space, *Lecture Notes in Computer Science* 1973 (2001) 420–434.
- [2] O. Aristondo, J. García-Lapresta, C. Lasso de la Vega, R. Marques Pereira, Classical inequality indices, welfare and illfare functions, and the dual decomposition, *Fuzzy Sets and Systems* 228 (2013) 114–136.
- [3] G. Beliakov, S. James, Unifying approaches to consensus across different preference representations, *Applied Soft Computing* 35 (2015) 888–897.
- [4] G. Beliakov, S. James, D. Nimmo, Can indices of ecological evenness be used to measure consensus?, in: *Proc. IEEE Intl. Conf. Fuzzy Systems'15*, Beijing, China, 2014, pp. 1–8.
- [5] K. Beyer, J. Goldstein, R. Ramakrishnan, U. Shaft, When is nearest neighbor meaningful?, in: C. Beeri, P. Buneman (Eds.), *Proc. ICDT*, Springer-Verlag, 1998, pp. 217–235.

- [6] J.C. Bezdek, Pattern Recognition with Fuzzy Objective Function Algorithms, Springer, 1981.
- [7] C. Bonferroni, Elementi di statistica generale, Libreria Seber, Firenze, 1930.
- [8] S. Bortot, R. Marques Pereira, On a new poverty measure constructed from the exponential mean, in: Proc. IFSA/EUSFLAT'15, Gijón, Spain, 2015, pp. 333–340.
- [9] S. Brin, Near neighbor search in large metric spaces, in: In Proceedings of the 21th International Conference on Very Large Data Bases, Morgan Kaufmann Publishers, 1995, pp. 574–584.
- [10] R. Cai, Z. Zhang, A.K. Tung, C. Dai, Z. Hao, A general framework of hierarchical clustering and its applications, Information Sciences 272 (2014) 29–48.
- [11] J. Camargo, Must dominance increase with the number of subordinate species in competitive interactions?, Journal of Theoretical Biology 161 (1993) 537–542.
- [12] H. Chang, D. Yeung, Robust path-based spectral clustering, Pattern Recognition 41 (2008) 191–203.
- [13] E. Chavez, G. Navarro, R. Baeza-Yates, J.L. Marroquin, Searching in metric spaces, ACM Computing Surveys 33 (2001) 273–321.
- [14] S. Dasgupta, Performance guarantees for hierarchical clustering, in: Proceedings of the Conference on Learning Theory, 2002, pp. 351–363.
- [15] I. Dimitrovski, D. Kocev, S. Loskovska, S. Džeroski, Improving bag-of-visual-words image retrieval with predictive clustering trees, Information Sciences 329 (2016) 851–865.
- [16] L.P. Dinu, R.T. Ionescu, Clustering methods based on closest string via rank distance, in: 14th Intl. Symp. Symbolic and Numeric Algorithms for Scientific Computing, IEEE, 2012, pp. 207–213.
- [17] D. Eddelbuettel, Seamless R and C++ Integration with Rcpp, Springer, New York, 2013.
- [18] L.N. Ferreira, L. Zhao, Time series clustering via community detection in networks, Information Sciences 326 (2016) 227–242.
- [19] R. Fisher, The use of multiple measurements in taxonomic problems, Annals of Eugenics 7 (1936) 179–188.

- [20] E. Fowlkes, C. Mallows, A method for comparing two hierarchical clusterings, *Journal of the American Statistical Association* 78 (1983) 553–569.
- [21] P. Fränti, O. Virtajoki, Iterative shrinking method for clustering problems, *Pattern Recognition* 39 (2006) 761–765.
- [22] L. Fu, E. Medico, FLAME, a novel fuzzy clustering method for the analysis of DNA microarray data, *BMC bioinformatics* 8 (2007) 3.
- [23] M. Gagolewski, Spread measures and their relation to aggregation functions, *European Journal of Operational Research* 241 (2015) 469–477.
- [24] J. García-Lapresta, C. Lasso de la Vega, R. Marques Pereira, A. Urrutia, A new class of fuzzy poverty measures, in: *Proc. of IFSA/EUSFLAT2015*, Gijón, Spain, 2015, pp. 1140–1146.
- [25] C. Gini, *Variabilità e mutabilità*, C. Cuppini, Bologna, 1912.
- [26] A. Gionis, H. Mannila, P. Tsaparas, Clustering aggregation, *ACM Transactions on Knowledge Discovery from Data* 1 (2007) 4.
- [27] J. Gower, G. Ross, Minimum spanning trees and single linkage cluster analysis, *Journal of the Royal Statistical Society. Series C (Applied Statistics)* 18 (1969) 54–64.
- [28] R. Graham, P. Hell, On the history of the minimum spanning tree problem, *Annals of the History of Computing* 7 (1985) 43–57.
- [29] D. Gómez, E. Zarrazola, J. Yáñez, J. Montero, A divide-and-link algorithm for hierarchical clustering in networks, *Information Sciences* 316 (2015) 308–328.
- [30] Z. Halim, M. Waqas, S.F. Hussain, Clustering large probabilistic graphs using multi-population evolutionary algorithm, *Information Sciences* 317 (2015) 78–95.
- [31] T. Hastie, R. Tibshirani, J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, Springer, 2013.
- [32] C. Heip, A new index measuring evenness, *Journal of Marine Biological Association of the United Kingdom* 54 (1974) 555–557.
- [33] A. Jain, M. Law, Data clustering: A user’s dilemma, *Lecture Notes in Computer Science* 3776 (2005) 1–10.

- [34] F. Jiang, G. Liu, J. Du, Y. Sui, Initialization of k-modes clustering using outlier detection techniques, *Information Sciences* 332 (2016) 167–183.
- [35] M. Kobus, Attribute decomposition of multidimensional inequality indices, *Economics Letters* 117 (2012) 189–191.
- [36] M. Kobus, P. Miłoś, Inequality decomposition by population subgroups for ordinal data, *Journal of Health Economics* 31 (2012) 15–21.
- [37] J.B. Kruskal, On the shortest spanning subtree of a graph and the traveling salesman problem, *Proceedings of the American Mathematical Society* 7 (1956) 48–50.
- [38] I. Kärkkäinen, P. Fränti, Dynamic local search algorithm for the clustering problem, in: *Proc. 16th Intl. Conf. Pattern Recognition’02*, volume 2, IEEE, 2002, pp. 240–243.
- [39] P. Legendre, L. Legendre, *Numerical Ecology*, Elsevier Science BV, Amsterdam, 2003.
- [40] J.B. MacQueen, Some methods for classification and analysis of multivariate observations, in: *Proc. Fifth Berkeley Symp. on Math. Statist. and Prob.*, volume 1, University of California Press, Berkeley, 1967, pp. 281–297.
- [41] D. Müllner, Modern hierarchical, agglomerative clustering algorithms, *ArXiv:1109.2378 [stat.ML]* (2011).
- [42] D. Müllner, fastcluster: Fast hierarchical, agglomerative clustering routines for R and Python, *Journal of Statistical Software* 53 (2013) 1–18.
- [43] F. Murtagh, A survey of recent advances in hierarchical clustering algorithms, *The Computer Journal* 26 (1983) 354–359.
- [44] F. Murtagh, P. Legendre, Ward’s hierarchical agglomerative clustering method: Which algorithms implement ward’s criterion?, *Journal of Classification* 31 (2014) 274–295.
- [45] C.F. Olson, Parallel algorithms for hierarchical clustering, *Parallel Computing* 21 (1995) 1313–1325.
- [46] W. Pedrycz, Conditional fuzzy c-means, *Pattern Recognition Letters* 17 (1996) 625–631.

- [47] W. Pedrycz, A. Bargiela, Granular clustering: A granular signature of data, *IEEE Transactions on Systems, Man, and Cybernetics, Part B: Cybernetics* 32 (2002) 212–224.
- [48] W. Pedrycz, J. Waletzky, Fuzzy clustering with partial supervision, *IEEE Transactions on Systems, Man, and Cybernetics, Part B: Cybernetics* 27 (1997) 787–795.
- [49] E. Pielou, *An Introduction to Mathematical Ecology*, Wiley-Interscience, New York, 1969.
- [50] E. Pielou, *Ecological Diversity*, Wiley, New York, 1975.
- [51] R. Prim, Shortest connection networks and some generalizations, *Bell System Technical Journal* 36 (1957) 1389–1401.
- [52] R Development Core Team, *R: A language and environment for statistical computing*, R Foundation for Statistical Computing, Vienna, Austria, 2015. <http://www.R-project.org>.
- [53] M. Radavanovic, A. Nanopoulos, M. Ivanovic, Hubs in space: Popular nearest neighbors in high-dimensional data, *Journal of Machine Learning Research* 11 (2010) 2487–2531.
- [54] F. Rohlf, Hierarchical clustering using the minimum spanning tree, *The Computer Journal* 16 (1973) 93–95.
- [55] C. Veenman, M. Reinders, E. Backer, A maximum variance cluster algorithm, *IEEE Transactions on Pattern Analysis and Machine Intelligence* 24 (2002) 1273–1280.
- [56] R. Xu, D.C. Wunsch II, *Clustering*, Wiley-IEEE Press, 2009.
- [57] P.N. Yianilos, Data structures and algorithms for nearest neighbor search in general metric spaces, in: *Proceedings of the Fourth Annual ACM-SIAM Symposium on Discrete Algorithms, SODA '93*, Society for Industrial and Applied Mathematics, 1993, pp. 311–321.
- [58] C. Zahn, Graph-theoretical methods for detecting and describing gestalt clusters, *IEEE Transactions on Computers* C-20 (1971) 68–86.
- [59] S. Zahra, M.A. Ghazanfar, A. Khalid, M.A. Azam, U. Naeem, A. Prugel-Bennett, Novel centroid selection approaches for kmeans-clustering based recommender systems, *Information Sciences* 320 (2015) 156–189.

- [60] T. Zhang, R. Ramakrishnan, M. Livny, BIRCH: An efficient data clustering method for very large databases, in: Proc. ACM SIGMOD'96 Intl. Conf. Management of Data, ACM, 1996, pp. 103–114.